

Kubernetes Course Syllabus

Archer Infotech, Kothrud Pune | Classroom & live-online batches | Updated September 2026

Institute: Archer Infotech, Kothrud, Pune - IT training since 2009

Website: <https://archerinfotech.in>

Contact: +91 9850 678451 | info@archerinfotech.in

Mode: Classroom at Kothrud, Pune and live-online batches

Duration: 2 Months

Level: Advanced

Course page: <https://archerinfotech.in/courses/cloud-devops/kubernetes-training-in-pune>

Course Overview

Kubernetes is the dominant container orchestrator in Pune product engineering and a core requirement on most senior DevOps / SRE / Platform Engineer roles — Persistent Systems, BMC Software, Bajaj Finserv, Synechron, BMW TechWorks India, Mercedes-Benz R&D India, Mastercard Pune Tech Hub, plus most fintech and SaaS startups run their primary production workloads on it. Archer Infotech's Kubernetes training in Pune teaches the platform as it is actually used in 2026 — Kubernetes 1.30+ (extending toward 1.32 by mid-year), workloads (Deployments, StatefulSets, DaemonSets, Jobs), Services and Ingress (NGINX + Gateway API), Helm 3 and Kustomize, RBAC and Network Policies (Calico / Cilium), persistent storage, observability with Prometheus + Grafana + Loki + OpenTelemetry, plus GitOps with Argo CD. The course doubles as preparation for the Certified Kubernetes Administrator (CKA) exam. Classroom in Kothrud, online live, and weekend batches available.

Why Learn Kubernetes in 2026

Kubernetes is no longer optional for Pune DevOps / SRE / Platform Engineer hiring — it is the single biggest separator between junior and mid-level candidates on every shortlist. Indeed Pune lists more than 700 active Kubernetes-specific roles as of May 2026, with another 1,000+ DevOps / SRE / Platform JDs that list Kubernetes as a hard requirement. The biggest employers are Persistent Systems, BMC Software, Bajaj Finserv, Synechron, BMW TechWorks India, Mercedes-Benz R&D India, Mastercard Pune Tech Hub, Cummins, John Deere ETC, plus the IT services majors with platform-engineering practices. Senior Kubernetes engineers in Pune regularly earn 1.5–2x equivalent-experience pure-developer salaries because the role bundles deep operational expertise with on-call accountability.

What changed in 2026: Kubernetes 1.30+ is the production default, with Sidecar Containers now stable, ValidatingAdmissionPolicy (CEL-based admission control replacing many webhook use cases) GA, and improved native support for AI / GPU workloads via Dynamic Resource Allocation (DRA). Gateway API has eclipsed Ingress as the recommended traffic-routing API for new clusters. eBPF-based networking (Cilium) has become the default CNI in many production deployments, replacing Calico for new clusters. GitOps with Argo CD has become the deployment pattern of choice in Pune product engineering, replacing manual `kubectl apply` workflows. The CKA exam was refreshed to weight troubleshooting and security questions more heavily, with the 2024 changes largely baked in by 2026.

What this means for hiring: 2026 Pune Kubernetes JDs expect K8s 1.28+ at depth (architecture, workloads, services, ConfigMaps / Secrets, RBAC, Network Policies), Helm fluency, basic Argo CD GitOps, one observability stack (Prometheus + Grafana + Loki), troubleshooting at the `kubectl describe / logs / exec / debug` level, and at least one production deployment story. Senior roles add Cilium / eBPF, service mesh basics (Istio / Linkerd), Pod Security Standards, and FinOps via Kubecost. Archer Infotech's curriculum is rebuilt around exactly these expectations — engineering-first, troubleshooting-heavy, certification-ready.

- 700+ active Kubernetes-specific roles on Indeed Pune (May 2026)
- Another 1,000+ DevOps / SRE / Platform JDs list K8s as a hard requirement

- Kubernetes 1.30+ — Sidecars stable, Gateway API, DRA for GPUs
- Argo CD GitOps + Cilium eBPF — the modern Pune production pattern
- Senior K8s engineers earn 1.5–2x equivalent-experience pure devs
- Certification path — CKA (Certified Kubernetes Administrator) included in curriculum

Who This Course Is For

This course fits you if

- Working developer wanting to add Kubernetes for senior backend / DevOps / Platform roles
- Working DevOps engineer wanting to deepen K8s and add CKA to your resume
- System / Linux administrator transitioning into modern Cloud / Container Platform engineering
- Working Docker user (production or hobby) wanting to graduate to orchestration
- Senior engineer at a Pune company that is migrating to Kubernetes — you need to lead the migration credibly
- DevOps practitioner targeting CKA / CKAD / CKS certification path with proper hands-on prep

This course is not the right fit if

- If you have no Linux command-line comfort — Linux fluency is required from day 1; we run `kubectl` + Bash + YAML constantly
- If you have no Docker experience — take our Docker course first; Kubernetes assumes container fluency from day 1
- If you have no programming or scripting background — at least basic Bash / Python is required for the operators / controllers / Helm chart work
- If you cannot put in 10–12 hours per week of lab work outside class — Kubernetes is the most lab-heavy of all our tracks
- If you only want a CKA certificate sticker with no production-engineering depth — there are cheaper exam-only courses; ours is engineering-first
- If you have 4+ years of production K8s with Helm + Argo + service mesh — you'll be under-stretched; talk to us about advanced specialisations (CKS, multi-cluster, service mesh deep dive)

Detailed Curriculum

Module 1: Kubernetes Architecture & First Cluster (Week 1)

Kubernetes explained as a set of cooperating control loops rather than a list of commands, because almost every troubleshooting question later resolves to knowing which component owns a behaviour. API server, etcd, scheduler and controller manager on the control plane; kubelet, kube-proxy and containerd on each node.

The declarative model — desired state, reconciliation, and why `kubectl apply` is not `kubectl run` — is established this week. Every student ends with a local cluster on kind or Minikube, a working kubeconfig with multiple contexts, and a first Deployment, Service and Ingress serving traffic.

- Control plane and data plane responsibilities
- API server, etcd, scheduler, controller manager
- kubelet, kube-proxy and containerd
- The declarative model and reconciliation loops
- `kubectl` essentials, kubeconfig and contexts
- kind and Minikube for local clusters
- Managed clusters — EKS, AKS, GKE differences
- First workload — Deployment, Service and Ingress

Module 2: Pods, Probes & the Container Lifecycle (Week 2)

The Pod as the unit of scheduling, and everything that happens inside one. Single and multi-container Pods, init containers for ordered setup, sidecar containers (now a stable, first-class feature) and ephemeral containers for debugging a running Pod without restarting it.

Probes get an extended treatment because misconfigured probes cause more self-inflicted outages than almost anything else: readiness versus liveness versus startup, timing parameters, and the failure modes each one produces when set wrongly. Resource requests, limits and QoS classes close the module.

- Pod anatomy, shared network and volume namespaces
- Multi-container patterns — sidecar, ambassador, adapter
- Init containers and ordered startup
- Sidecar containers as a stable feature
- Ephemeral containers for live debugging
- Readiness, liveness and startup probes
- Probe timing and the outages bad settings cause
- Resource requests, limits and QoS classes
- Pod lifecycle, restart policies and termination grace

Module 3: Deployments, StatefulSets, DaemonSets & Jobs (Week 2)

The workload controllers and what each one guarantees. Deployments and ReplicaSets with rolling-update parameters, revision history, rollback, and pause and resume for staged releases; the ownership chain that explains why deleting a ReplicaSet does not do what people expect.

StatefulSets cover stable network identity, ordered deployment and scaling, and per-replica storage. DaemonSets handle node-level agents, and Jobs and CronJobs cover batch work including parallelism, completions, backoff limits and the concurrency policy.

- Deployments, ReplicaSets and ownership chains
- Rolling updates — maxSurge and maxUnavailable
- Revision history, rollback, pause and resume
- Recreate vs RollingUpdate strategies
- StatefulSets — stable identity and ordered operations
- Per-replica volume claim templates
- DaemonSets for node-level agents
- Jobs — parallelism, completions, backoff limits
- CronJobs and concurrency policy

Module 4: Services, Endpoints & Cluster DNS (Week 3)

How traffic finds a Pod. The four Service types and what each actually creates, EndpointSlices and how they track healthy backends, and the kube-proxy modes (iptables and IPVS) that implement the routing underneath.

CoreDNS and the cluster DNS naming scheme come next, including headless Services for direct Pod addressing and the resolution behaviour StatefulSets depend on. The module is deliberately heavy on debugging: `kubectl describe`, endpoint inspection, DNS lookups from inside a Pod, and port-forward.

- Service types — ClusterIP, NodePort, LoadBalancer, ExternalName
- Endpoints, EndpointSlices and readiness gating
- kube-proxy modes — iptables and IPVS
- CoreDNS and the cluster DNS naming scheme
- Headless Services and direct Pod addressing
- Service topology and internal traffic policy
- Debugging — describe, endpoints, nslookup, port-forward
- Common causes of a Service returning no backends

Module 5: Ingress, Gateway API & TLS (Week 3)

External access done properly. The Ingress resource and the NGINX Ingress Controller — path and host routing, annotations, rewrite rules, and TLS termination with cert-manager issuing certificates from Let's Encrypt automatically.

Gateway API is taught as the successor rather than an alternative: Gateway, HTTPRoute, GRPCRoute, the role separation between infrastructure and application teams that motivated it, and a migration path from Ingress. Header-based routing and traffic splitting close the module.

- Ingress resources, classes and controllers
- NGINX Ingress Controller — routing and annotations
- TLS termination and SNI

- cert-manager and automated certificate issuance
- Gateway API — Gateway, HTTPRoute, GRPCRoute
- Role separation between platform and app teams
- Header-based routing and traffic splitting
- Migrating from Ingress to Gateway API

Module 6: Configuration, Secrets & External Secret Stores (Week 4)

Separating configuration from images. ConfigMaps as environment variables, arguments and mounted files; immutable ConfigMaps and Secrets for performance and safety; and the reload problem — what actually happens to a running Pod when a mounted ConfigMap changes.

Secrets are covered honestly: base64 is encoding, not encryption, so the module moves to encryption at rest, then to the External Secrets Operator backed by AWS Secrets Manager, Azure Key Vault or Vault, and Sealed Secrets for storing encrypted secrets safely in a GitOps repository.

- ConfigMaps as env vars, args and mounted files
- Immutable ConfigMaps and Secrets
- Config reload behaviour and restart strategies
- Secrets and the limits of base64
- Encryption at rest for etcd
- External Secrets Operator with cloud secret stores
- Sealed Secrets for GitOps repositories
- Downward API and pod metadata injection

Module 7: Persistent Storage, CSI & Stateful Workloads (Week 4)

Running stateful software on Kubernetes, and the judgement of when not to. PersistentVolumes and claims, access modes, storage classes and dynamic provisioning, reclaim policies, and the CSI driver model with cloud implementations — EBS, Azure Disk, Persistent Disk — plus Longhorn for on-premises.

Volume snapshots, restore and expansion follow. The module ends with real StatefulSet patterns for PostgreSQL, Redis and Kafka, and a frank discussion of when a managed database is the better engineering decision.

- PersistentVolumes, claims and access modes
- Storage classes, dynamic provisioning, reclaim policies
- CSI drivers — EBS, Azure Disk, Persistent Disk, Longhorn
- Volume snapshots, restore and expansion
- StatefulSet volume claim templates
- Postgres, Redis and Kafka patterns on Kubernetes
- Backup strategy for stateful workloads
- When a managed database beats running your own

Module 8: Helm 3 & Kustomize (Week 5)

Packaging and templating manifests. Helm 3 charts end to end — structure, values files and precedence, the template language and its common traps, named templates, hooks, chart dependencies, and release lifecycle including upgrade, rollback and history. OCI registries as the modern chart distribution path.

Kustomize follows with bases, overlays, patches and generators — a genuinely different philosophy, not a competitor with the same shape. The module ends with a clear-eyed comparison and guidance on choosing per team rather than per preference.

- Helm chart structure and the template language
- Values files, precedence and schema validation
- Named templates, helpers and common traps
- Hooks and the release lifecycle
- Chart dependencies and subchart values
- OCI registry support for chart distribution

- Kustomize bases, overlays, patches and generators
- Helm vs Kustomize — choosing per team
- Deploying kube-prometheus-stack via Helm

Module 9: CRDs, Controllers & the Operator Pattern (Week 5)

How Kubernetes extends itself, which is what separates an operator of clusters from an engineer who can shape one. Custom Resource Definitions, schema validation with OpenAPI, versioning and conversion, and how a CRD plus a controller reproduces the same reconciliation model the built-in resources use.

The operator pattern is then read in practice — examining a real operator's reconcile loop, its status conditions and finalisers — with an overview of the Operator SDK and Kubebuilder for teams that need to write their own.

- Custom Resource Definitions and OpenAPI schemas
- CRD versioning and conversion webhooks
- The controller reconciliation loop
- Status conditions, finalisers and owner references
- Reading a real operator's source
- Operator SDK and Kubebuilder overview
- Operator lifecycle and when to adopt one
- Aggregated API servers in brief

Module 10: RBAC, Service Accounts & Admission Control (Week 6)

The authorisation and admission layer. Roles and ClusterRoles, RoleBindings and ClusterRoleBindings, the aggregation model, and auditing what a given principal can actually do with `kubectl auth can-i`. Service accounts, projected and bound tokens, and cloud workload identity for keyless access to external resources.

Admission control closes the module: the mutating and validating webhook flow, and ValidatingAdmissionPolicy with CEL expressions — the in-tree, webhook-free path that is now the default recommendation for straightforward policy.

- Roles, ClusterRoles and the two binding kinds
- Role aggregation and built-in roles
- Auditing access with `kubectl auth can-i`
- Service accounts, projected and bound tokens
- Cloud workload identity for keyless access
- The admission chain — mutating then validating
- ValidatingAdmissionPolicy with CEL
- Admission webhooks and their failure modes

Module 11: Network Policies, CNI & Pod Security (Week 6)

Hardening the cluster's runtime. Network Policies for ingress and egress control, the default-allow behaviour that surprises people, and building a default-deny baseline that does not break DNS. Calico versus Cilium as CNI choices, and eBPF-based networking and observability with Cilium and Hubble.

Pod Security Standards and Pod Security Admission replace the removed PodSecurityPolicy, and the module ends on image security — Trivy scanning, Cosign signatures and admission-time verification.

- Network Policies — ingress, egress, default-allow
- Building a default-deny baseline without breaking DNS
- Calico vs Cilium — choosing the CNI
- eBPF networking and Hubble observability
- Pod Security Standards — privileged, baseline, restricted
- Pod Security Admission enforcement modes
- Security context, capabilities and read-only root
- Image scanning with Trivy
- Cosign signatures and admission-time verification

Module 12: Observability — Prometheus, Grafana, Loki & Tempo (Week 7)

Seeing inside a running cluster. Prometheus architecture and the pull model, exporters, and the Prometheus Operator's ServiceMonitor and PodMonitor resources. PromQL taught as a language — selectors, rate versus increase, histograms and quantiles, recording rules and query cost.

Alertmanager covers routing, grouping, inhibition and silences; Grafana dashboards are built to the USE, RED and Four Golden Signals methods. Loki with LogQL, Tempo for traces and OpenTelemetry instrumentation complete the stack, ending in an on-call simulation against a deliberately broken cluster.

- Prometheus architecture, storage and the pull model
- Exporters, ServiceMonitors and PodMonitors
- PromQL — selectors, operators, functions
- rate vs increase, histograms and quantiles
- Recording rules and query cost
- Alertmanager — routing, grouping, inhibition, silences
- Grafana dashboards — USE, RED, Four Golden Signals
- Loki and LogQL, with label-cardinality discipline
- Tempo, distributed tracing and OpenTelemetry
- On-call simulation — debugging a controlled outage

Module 13: Scheduling, Autoscaling & Cluster Operations (Week 8)

Deciding where Pods run and how a cluster grows. The scheduler's filter and score phases, node selectors and affinity, pod affinity and anti-affinity for spreading replicas, taints and tolerations, topology spread constraints, and pod priority with preemption.

Autoscaling covers all three axes — Horizontal Pod Autoscaler including custom metrics, Vertical Pod Autoscaler, and node-level scaling with Cluster Autoscaler and Karpenter. Cluster operations closes the module: Pod Disruption Budgets, draining, upgrade strategy and etcd backup and restore.

- Scheduler filter and score phases
- Node selectors, node affinity and anti-affinity
- Pod affinity and topology spread constraints
- Taints, tolerations and dedicated node pools
- Pod priority, preemption and eviction
- Horizontal Pod Autoscaler with custom metrics
- Vertical Pod Autoscaler and its limits
- Cluster Autoscaler and Karpenter
- Pod Disruption Budgets, draining and upgrades
- etcd backup and restore

Module 14: GitOps with Argo CD & Service Mesh (Week 8)

The delivery model most Kubernetes-first teams have settled on. Argo CD architecture, Application CRDs, sync policies with self-healing and pruning, sync waves and health assessment, and ApplicationSets for generating applications across clusters and environments. Repository structure and promotion from dev to staging to production are built as code rather than as a runbook.

Service mesh is covered comparatively — Istio versus Linkerd, the sidecar and ambient models, mTLS, traffic shifting and observability — with honest guidance on when a mesh is not yet worth its operational cost.

- GitOps principles and reconciliation
- Argo CD architecture and Application CRDs
- Sync policies, self-healing, pruning and sync waves
- ApplicationSets for multi-cluster generation
- Repository structure and environment promotion
- Secrets in GitOps — Sealed Secrets and External Secrets
- Service mesh — Istio vs Linkerd, sidecar vs ambient
- mTLS, traffic shifting and mesh observability

- When a service mesh is not worth the cost

Module 15: Capstone Project & CKA Exam Preparation (Weeks 9–10 + 1 week placement prep)

Full-time capstone work plus structured Certified Kubernetes Administrator preparation. The capstone is a complete platform — Helm-packaged workloads, Argo CD delivery, the full observability stack, network policy and RBAC baselines, and documented runbooks.

CKA preparation is hands-on rather than multiple choice, because the exam is: timed troubleshooting drills, `kubectl` shortcuts and aliases that buy back minutes, and two full-length mock exams. Interview preparation covers a cluster-troubleshooting round and an architecture round, plus resume, LinkedIn and portfolio polish.

- Capstone implementation, deployment and README
- Runbooks and architecture decision records
- CKA-style hands-on troubleshooting drills
- Two full-length timed mock CKA exams
- `kubectl` shortcuts, aliases and imperative commands
- Cluster troubleshooting mock interview round
- Kubernetes architecture and scenario round
- Resume and LinkedIn rewrite for K8s / Platform JDs
- GitHub portfolio polish — Helm charts, Argo apps, dashboards
- HR mock interview and salary negotiation

Module 16: AI-Assisted Development Workflow (Final Week)

The skill every 2026 hiring panel now probes for — building real work with AI in the loop, responsibly. Learn to drive AI assistants (GitHub Copilot, Claude, Cursor, and IDE-native AI) to scaffold and accelerate the tools and stack this course covers, generate tests, explain and refactor unfamiliar code, and cut the boilerplate — while keeping you firmly in control of every decision. Heavy focus on guardrails: reviewing each AI suggestion, spotting hallucinated APIs or wrong answers, and handling licensing and data-privacy concerns. Close with a mini-project that takes a deliverable end-to-end using an AI-assisted workflow, then fold the same tooling into version control and everyday team practice.

- AI assistants — GitHub Copilot, Claude, Cursor, IDE-native AI
- Effective prompting for this course's stack — scaffolding, boilerplate, config
- AI-assisted test generation and coverage
- Explaining, refactoring, and modernising unfamiliar code with AI
- AI-driven review, error detection, and quality checks
- Generating and maintaining documentation with AI
- AI debugging — interpreting errors, logs, and failing output
- Guardrails — reviewing output, avoiding hallucinations, licensing & data privacy
- Team workflow — AI in the editor, in reviews, and in delivery pipelines
- Mini-project — a deliverable built end-to-end with an AI-assisted workflow

Capstone Projects

Production-Grade Cluster with kube-prometheus-stack + Argo CD

A complete production-style Kubernetes cluster on EKS / AKS — cluster provisioned via Terraform, Argo CD installed via Helm and managing itself (the App-of-Apps pattern), three sample workloads (web frontend, REST API, worker) deployed via Helm charts in a Git repository, kube-prometheus-stack for metrics, Loki for logs, Sealed Secrets / External Secrets for credential management, NGINX Ingress with cert-manager for TLS, plus a small Network Policy baseline. The full cluster stands up via Terraform + Argo CD bootstrap and tears down clean. Outcome: a public GitHub repository with the cluster definition, Helm charts, and Grafana dashboards exported as JSON.

- Technology: Terraform 1.7+
- Technology: EKS or AKS
- Technology: Argo CD with App-of-Apps
- Technology: Helm 3 charts

- Technology: kube-prometheus-stack
- Technology: Loki + Tempo
- Technology: Sealed Secrets / External Secrets
- Technology: cert-manager + NGINX Ingress

Multi-Tenant Platform with Network Policies + Cilium

A multi-tenant Kubernetes platform — three tenant namespaces with strict Network Policy isolation (deny-by-default ingress + egress), per-tenant RBAC and ResourceQuotas, Cilium as the CNI for eBPF observability (Hubble for live traffic visualisation), Pod Security Standards enforced, and a small admin Operator built with kubebuilder that manages tenant lifecycle. Demonstrates the patterns Pune Platform Engineer roles hire on — multi-tenant, secure, observable, extensible.

- Technology: Cilium with Hubble
- Technology: Network Policies (deny-by-default)
- Technology: Pod Security Standards
- Technology: RBAC + ResourceQuotas
- Technology: Operator built with kubebuilder
- Technology: External Secrets Operator

GitOps Multi-Environment Promotion with Argo CD

A complete GitOps deployment platform — Argo CD with ApplicationSets managing dev / staging / prod environments, Helm charts with environment-specific values, automated promotion via PR-driven flow (PR for dev, tag for staging, manual approval for prod), Sealed Secrets for credential management in Git, plus a Slack notification integration on sync events. Demonstrates the deployment pattern of choice in Pune product engineering and is the artefact that opens senior Platform Engineer interviews.

- Technology: Argo CD + ApplicationSets
- Technology: Helm 3 with multi-env values
- Technology: Sealed Secrets
- Technology: GitHub Actions for PR flow
- Technology: Slack notifications
- Technology: Optional: Argo Rollouts for canary

Career Outcomes

Kubernetes Administrator, Platform Engineer, and SRE are among the highest-paid technical roles in Pune in 2026 — Indeed Pune lists 700+ active Kubernetes-specific openings, plus another 1,000+ DevOps / SRE / Platform JDs that list Kubernetes as a hard requirement. Senior compensation regularly exceeds equivalent-experience full-stack developer offers because the role bundles deep operational expertise with on-call accountability. The biggest Pune employers are Persistent Systems, BMC Software, Bajaj Finserv, Synechron, BMW TechWorks India, Mercedes-Benz R&D India, Mastercard Pune Tech Hub, Cummins, John Deere ETC, plus the IT services majors with platform-engineering practices.

What pulls a Kubernetes engineer above the median band: a public GitHub repository with a real Argo CD GitOps deployment, demonstrable Helm chart authoring, one production observability stack you can defend, troubleshooting depth (the CKA-style hands-on skills), and the CKA certificate. Our capstone projects and certification track are designed exactly around these signals.

Senior Platform Engineer / SRE bands at the top end are reported as national figures (Pune-specific Indeed pages do not exist for those roles); Pune trends within $\pm 10\%$ of these figures based on AmbitionBox and 6figr.

Roles after this course

- Kubernetes Administrator
- Site Reliability Engineer (SRE)
- Platform Engineer
- DevOps Engineer (Kubernetes-focused)
- Cloud-Native Engineer
- Junior Cluster Operator

- Container Platform Engineer

Indicative salary bands

- Kubernetes Administrator (Pune) - Rs. 8,40,000 per year average (source: Indeed Pune (Kubernetes Administrator))
- Junior Kubernetes Engineer (Pune entry, <2 years) - Rs. 5,00,000 – Rs. 8,00,000 per year (source: AmbitionBox Pune Kubernetes Engineer)
- Mid-level Platform / DevOps Engineer with K8s (Pune, 3–5 years) - Rs. 14,00,000 – Rs. 22,00,000 per year (source: Glassdoor Pune Platform Engineer)
- Senior Site Reliability Engineer (Pune, 5–8 years) - Rs. 18,00,000 – Rs. 30,00,000 per year (source: Glassdoor Pune SRE)
- Lead / Staff Platform Engineer (national, 8+ years) - Rs. 30,00,000 – Rs. 55,00,000 per year (source: 6figr India Lead Platform Engineer (Pune ±10%))

Companies hiring for these skills

- Persistent Systems
- BMC Software
- Bajaj Finserv
- Synechron
- BMW TechWorks India
- Mercedes-Benz R&D India
- Mastercard Pune Tech Hub
- Cummins India
- John Deere ETC
- Honeywell
- TCS
- Infosys
- Cognizant
- Capgemini
- Atos / Eviden
- Mphasis

Batch Modes and Duration

Duration: 8 weeks of structured curriculum plus 2 weeks of capstone project work and CKA / interview preparation (~2.5 months total)

Classroom

- Location: Archer Infotech, Kothrud, Pune
- Morning batch — 10:00 to 13:00
- Evening batch — 18:00 to 21:00
- Lab access available outside class hours

Live online

- Same hours as classroom batches — morning or evening
- Recordings available for review
- Same lab reviews and project feedback as in-person batches
- Zoom for live sessions
- Personal AWS / Azure sandbox for managed-cluster labs
- GitHub for code, Helm chart, and Argo CD manifest reviews
- Slack / WhatsApp for asynchronous Q&A

Weekend

- Saturday + Sunday, 09:00 to 13:00
- Stretches over ~4 months instead of 2.5 to accommodate working professionals. Same content, lower weekly load.

Maximum 15 students per batch — small enough that the trainer reviews every student's manifests and Helm charts personally. Classroom batches start every 4 weeks; weekend batches every 6 weeks.

Fees

Course fees range from Rs. 20,000 to Rs. 90,000 depending on mode (classroom / online / weekend), batch type, and any applicable concession. Kindly reach us for the current 2026 quote — we calibrate by early-bird timing, group enrolment, and returning-alumni concessions. The CKA exam voucher (USD ~395 / ~Rs. 33,000, often discounted via CNCF promotions) is paid directly by the student and is not part of our fee. Cloud sandbox spend (managed-cluster labs) typically runs Rs. 500–Rs. 1,500 across the course.

Indicative range: Rs. 20,000 – Rs. 90,000 — the higher end covers placement-track classroom batches with full CKA mock-exam track, multi-tenant + GitOps capstones, and extended interview prep; the lower end covers concession-eligible online or weekend formats.

- Single payment with early-bird discount
- EMI in 2–3 instalments at no extra cost
- Corporate sponsorship — invoiced to your employer with GST

Placement Support

Placement support starts from week 6 of the course, not at the end. By the time you finish the curriculum, your resume highlights real Argo CD GitOps work, your GitHub has a deployable cluster + Helm + Grafana reference repository, and you have completed at least three mock technical interviews against question banks from Pune Kubernetes / Platform hiring teams.

We say placement support, not placement guarantee — for two honest reasons. First, no institute can guarantee a hire when the final decision is the company's. Second, the institutes that do guarantee tend to bury the conditions in fine print. Our support is unconditional, time-bound (six months after course completion), and includes free re-entry to a future batch's interview-prep sessions if your first round of interviews does not land.

The placement process

- Week 6 — resume and LinkedIn rewrite, calibrated for K8s / Platform / SRE JDs
- Week 7 — GitHub portfolio cleanup, Helm chart polish, Grafana dashboard exports
- Weeks 8–9 — CKA-style troubleshooting drills, architecture mock rounds
- Weeks 9–10 — three rounds of mock technical interviews
- Week 10 — HR mock interview and salary negotiation coaching
- Post-course — referrals via our 17-year alumni network at 12+ partner companies
- Up to 6 months of continued support after course end
- Free re-entry to future batch interview-prep sessions if first round does not land

Partner companies

- Persistent Systems
- BMC Software
- Bajaj Finserv
- Synechron
- BMW TechWorks India
- Mercedes-Benz R&D India
- Mastercard Pune Tech Hub
- Cummins
- TCS
- Infosys
- Cognizant
- Capgemini

Prerequisites and How to Start

Prerequisites: Docker fluency (Dockerfile authoring, multi-stage builds, Compose at a working level — if you have not used Docker production-style, take our Docker course first), Linux command-line comfort, basic Bash / Python scripting, and willingness to commit 10–12 hours per week of lab work outside class. Working DevOps engineers, backend developers with Docker experience, and Linux administrators typically slot in well; pure Windows-administrator candidates with no container background should do Docker + Linux foundations first.

Suggested first steps

- Decide your mode — classroom in Kothrud, online live, or weekend
- Check the upcoming batch dates on our batch schedule page
- Book a free 30-minute counselling call — we will honestly tell you whether the course fits your goal (we say no to roughly 15% of K8s enquirers because Docker / Linux foundation isn't yet there)
- Confirm enrolment and complete pre-course orientation (Docker, kubectl, kind / Minikube install scripts; AWS / Azure account creation guide for managed-cluster labs)
- Show up to day one with a laptop running 64-bit Linux / macOS / Windows-with-WSL2, 16GB+ RAM (recommended for local clusters), and Docker Desktop or Docker Engine pre-installed

Frequently Asked Questions

Which is the best Kubernetes training institute in Pune?

We can't honestly answer 'best' for ourselves. The test that works: ask any institute you are considering to (1) name the trainer who will teach your batch and show their LinkedIn, (2) show real student Helm charts and Argo CD applications on GitHub, and (3) name companies that hired their last 5 batches. Compare on those three.

How long does Kubernetes training in Pune take at Archer Infotech?

Approximately 2.5 months — 8 weeks of structured curriculum plus 2 weeks of capstone project and CKA / interview preparation. The weekend batch stretches over ~4 months at the same content depth, designed for working professionals.

What is the salary of a Kubernetes Engineer in Pune?

Indeed Pune reports an average of Rs. 8.40 lakh per year for Kubernetes Administrator. Junior Kubernetes Engineer Pune entry sits at Rs. 5–8 lakh per year per AmbitionBox. Mid-level Platform / DevOps Engineers with K8s (3–5 years) earn Rs. 14–22 lakh per Glassdoor. Senior SREs (5–8 years) earn Rs. 18–30 lakh. Lead / Staff Platform Engineers earn Rs. 30–55 lakh nationally with Pune trending within ±10%.

What is the fee for the Kubernetes course in Pune?

Course fees range from Rs. 20,000 to Rs. 90,000 depending on mode (classroom / online / weekend), batch type, and applicable concession. The higher end covers placement-track classroom batches with full CKA mock-exam track and extended interview prep; the lower end covers concession-eligible online or weekend formats. Cloud sandbox spend across the course typically runs Rs. 500–Rs. 1,500 (paid directly to AWS / Azure). The CKA exam voucher (~Rs. 33,000) is paid directly to CNCF.

Does the course prepare me for the CKA certification?

Yes — CKA preparation is woven through the curriculum and concentrated in weeks 9–10. Two full-length timed mock exams plus CKA-style hands-on troubleshooting drills are part of the course. Most students who complete the lab work seriously pass the live exam on first attempt.

Should I learn Docker first?

Yes — Docker fluency is required from day 1. We do not start from 'what is a container'. If you have not used Docker production-style (Dockerfile authoring, multi-stage builds, Compose), take our Docker course first; that is 1.5 months and ramps you up to the level this course assumes.

Will I work on real projects?

Yes — three capstone projects: (1) production-grade cluster with kube-prometheus-stack + Argo CD on EKS / AKS, (2) multi-tenant platform with Network Policies + Cilium + custom Operator, (3) GitOps multi-environment promotion with Argo CD + ApplicationSets. All three become public GitHub repositories with cluster manifests, Helm charts, and Grafana dashboard exports.

Do I get my own cluster during the course?

Yes — every student runs a local kind / Minikube cluster from day 1 plus a managed cluster (EKS or AKS — student's choice) for production-realistic exercises. Cloud spend stays under Rs. 1,500 across the course if you tear down resources after labs (which we drill the discipline of).

Is Argo CD / GitOps covered?

Yes — week 8 is dedicated to Argo CD with ApplicationSets and the dev -> staging -> prod promotion pattern. Capstone Projects #1 and #3 use Argo CD as the deployment driver. GitOps has become the deployment standard in Pune product engineering; we treat it as core, not advanced.

What about service mesh — Istio / Linkerd?

Service mesh is covered at primer depth in week 8 — what it is, when it earns its complexity, what Istio and Linkerd each do. We do not deep-dive Istio in this course because it would deserve a separate 4-week track; service mesh interview questions at the primer level are what most Pune K8s panels ask, and we cover that fully.

Are weekend Kubernetes classes available in Pune?

Yes — Saturday and Sunday, 09:00–13:00, stretched over ~4 months instead of 2.5. Same content, same trainer, same labs and capstone. Designed for working professionals who cannot attend weekday batches.

How does this course compare to your DevOps course?

The DevOps course is the broader programme — Linux + Git + Docker + K8s + CI/CD + Terraform + Observability + Security + FinOps over 3 months. This Kubernetes course is the focused deep-dive — 2.5 months on Kubernetes specifically with CKA prep included. Pick the DevOps course if you are starting fresh and need the broad foundation; pick this course if you already have Linux / Docker / CI/CD covered and want K8s depth + CKA.

What support do I get after course completion?

Six months of active placement support — mock interviews calibrated for Kubernetes / SRE / Platform Engineer roles (troubleshooting + architecture + behavioural rounds), referrals via our alumni network at 12+ partner companies, resume / LinkedIn / GitHub rewrites, and salary negotiation coaching. If your first round of interviews does not land, you can sit in on a future batch's interview-prep sessions free of charge.

Is the named trainer actually teaching, or are they just on the brochure?

Amol Patil personally leads every session of every batch from Day 1 through capstone — he ships Kubernetes manifests and Helm charts daily and brings real production patterns into the classroom. The same name on this page is the same person you meet on day one; his LinkedIn is on the trainer profile page, and we welcome a 30-minute conversation with him before you enrol.

Enquire or book a free demo session. Archer Infotech, Flat No. 12, Divyadarshan Housing Society, behind Kothrud Bus Stand Road, Kothrud, Pune 411038. Phone +91 9850 678451. Email info@archerinfotech.in. Web archerinfotech.in.